

# Introduction to Robotics


## Week-2

Prof. Balaraman Ravindran

IIT-Madras



- ▶ Gaussian Filters are a family of recursive state estimators for continuous spaces, that all share the idea that beliefs are represented by multivariate normal distributions.



$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$

- ▶ This choice to represent beliefs by a Gaussian distribution creates a model where there is a small margin of uncertainty around a single true state. Gaussian distributions are a poor choice for estimation problems where many distinct hypothesis exist.

### The Kalman Filter

Besides the Markov assumption, the belief distribution estimated by the Kalman algorithm will be a Gaussian distribution if the following three properties hold:

1. The system must follow linear dynamics, with randomness in state transition modelled as Gaussian noise.

$$x_t = A_t x_{t-1} + B_t u_t + \epsilon_t.$$

$$\text{Here, } \epsilon_t \sim \mathcal{N}(0, R_t)$$

The state transition probability  $p(x_t | u_t, x_{t-1})$  is thus given by:

$$\frac{1}{(2\pi)^{n/2} |R_t|^{1/2}} \exp\left(-\frac{1}{2} (x_t - A_t x_{t-1} - B_t u_t)^T R_t^{-1} (x_t - A_t x_{t-1} - B_t u_t)\right)$$

$$\text{i.e. } \mathcal{N}(A_t x_{t-1} + B_t u_t, R_t)$$

- The measurement probability must also be linear in it's arguments, with added Gaussian noise:

$$z_t = C_t x_t + \delta_t$$

Here,  $\delta_t \sim \mathcal{N}(0, Q_t)$

The measurement probability  $p(z_t|x_t)$  is thus given by:

$$\frac{1}{(2\pi)^{n/2} |Q_t|^{1/2}} \exp\left(-\frac{1}{2}(x_t - C_t x_t)^T Q_t^{-1}(x_t - C_t x_t)\right)$$

i.e  $\mathcal{N}(C_t x_t, Q_t)$

- The initial belief  $bel(x_0)$  must be normal distributed with some mean  $\mu_0$  and covariance  $\Sigma_0$ . i.e  $\mathcal{N}(\mu_0, \Sigma_0)$


$$bel(x_0) = \frac{1}{(2\pi)^{n/2} |\Sigma_0|^{1/2}} \exp\left(-\frac{1}{2}(x_t - \mu_0)^T \Sigma_0^{-1}(x_t - \mu_0)\right)$$

- ▶ Similar to the **Bayes Filter** algorithm, the **Kalman Filter** algorithm recursively estimates  $bel(x_t)$  from  $\{bel(x_{t-1}), u_t, z_t\}$
- ▶  $bel(x_t)$  is parameterized by  $\mu_t, \Sigma_t$

```
1:  Algorithm Kalman_filter( $\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$ ):  
2:       $\bar{\mu}_t = A_t \mu_{t-1} + B_t u_t$   
3:       $\bar{\Sigma}_t = A_t \Sigma_{t-1} A_t^T + R_t$   
4:       $K_t = \bar{\Sigma}_t C_t^T (C_t \bar{\Sigma}_t C_t^T + Q_t)^{-1}$   
5:       $\mu_t = \bar{\mu}_t + K_t (z_t - C_t \bar{\mu}_t)$   
6:       $\Sigma_t = (I - K_t C_t) \bar{\Sigma}_t$   
7:      return  $\mu_t, \Sigma_t$ 
```

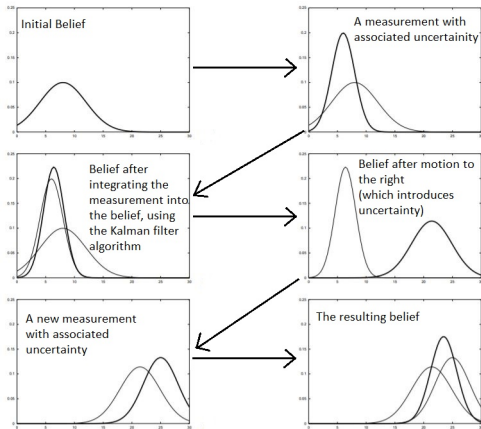
- ▶ **Note:** Similar to the Bayes filter, we have a **prediction** step (*Lines 2, 3*) - which modifies belief in accordance to an action, and a **measurement update** step (*Lines 4,5,6*) - in which sensor data is integrated into the present belief.
- ▶ The variable  $\mathbf{K}_t$ , computed in Line 4 is called *Kalman gain*. It specifies the degree to which the measurement is incorporated into the new state estimate.

### A simplistic one-dimensional localization scenario



- ▶ The robot moves along the horizontal axis.
- ▶ The robot queries it's GPS sensors on it's location


## An Illustration of the Kalman Filter Algorithm





### The Extended Kalman Filter

- ▶ The assumptions of linear state transitions and linear measurements with added Gaussian noise are rarely satisfied in practice. For example, a robot moving in a circular trajectory with constant velocity cannot be described by linear next state transitions.



- ▶ The extended Kalman filter (EKF) overcomes the linearity assumption, assuming instead that the next state probability and the measurement probabilities are governed by nonlinear functions  $g$  and  $h$ .

$$x_t = g(u_t, x_t) + \epsilon_t$$

$$z_t = h(x_t) + \delta_t$$

- ▶ Unfortunately, with arbitrary functions  $g$  and  $h$ , we are no longer guaranteed that the belief distribution estimated by the Kalman Filter algorithm is a Gaussian.
- ▶ The extended Kalman filter (EKF) overcomes this problem by calculating a linear approximation for the functions  $g$  and  $h$ , to approximate the true belief.
- ▶ Once  $g$  and  $h$  are linearized, the mechanics of belief propagation are equivalent to those of the Kalman filter

*Taylor expansion:*

Slope is given by the partial derivative:

$$g'(u_t, x_{t-1}) := \frac{\partial g(u_t, x_{t-1})}{\partial x_{t-1}}$$

and we approximate:

$$g(u_t, x_{t-1}) \approx g(u_t, \mu_{t-1}) + g'(u_t, \mu_{t-1})(x_{t-1} - \mu_{t-1})$$

Defining,  $G_t := g'(u_t, x_{t-1})$  (Jacobian - matrix of size  $n \times n$ )

We have,

$$g(u_t, x_{t-1}) \approx g(u_t, \mu_{t-1}) + G_t(x_{t-1} - \mu_{t-1})$$

So that next state probability  $p(x_t | u_t, x_{t-1})$  is approximated to:

$$\mathcal{N}(g(u_t, \mu_{t-1}) + G_t(x_{t-1} - \mu_{t-1}), R_t)$$

**Similarly,**

We approximate:

$$\begin{aligned} h(x_t) &\approx h(\bar{\mu}_t) + h'(\bar{\mu}_t)(x_t - \bar{\mu}_t) \\ &= h(\bar{\mu}_t) + H_t(x_t - \bar{\mu}_t) \end{aligned}$$

(Where  $\bar{\mu}_t$  is the state deemed most likely at the time of linearizing  $h$ )

So that measurement probability is approximated to:


$$\mathcal{N}(h(\bar{\mu}_t) + H_t(x_t - \bar{\mu}_t), Q_t)$$

- The linear predictions in Kalman filters are replaced by their nonlinear generalizations in EKF. Moreover, EKFs use Jacobians  $G_t$  and  $H_t$  instead of the corresponding linear system matrices  $A_t$ ,  $B_t$ , and  $C_t$  in Kalman filters.

```
1:  Algorithm Extended_Kalman_filter( $\mu_{t-1}, \Sigma_{t-1}, u_t, z_t$ ):  
2:       $\bar{\mu}_t = g(u_t, \mu_{t-1})$   
3:       $\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t$   
4:       $K_t = \bar{\Sigma}_t H_t^T (H_t \bar{\Sigma}_t H_t^T + Q_t)^{-1}$   
5:       $\mu_t = \bar{\mu}_t + K_t (z_t - h(\bar{\mu}_t))$   
6:       $\Sigma_t = (I - K_t H_t) \bar{\Sigma}_t$   
7:      return  $\mu_t, \Sigma_t$ 
```

## Non-Parametric Filters

- ▶ Alternative to Gaussian Filters, non-parametric filters do not assume a fixed functional form over the belief distribution, allowing us to model more complex, non-gaussian probability distributions.



- ▶ They estimate the belief distribution at a fixed number of positions in the state space.

### The Particle Filter

- ▶ Instead of representing the belief distribution by a parametric form, particle filters represent this distribution by a set of samples drawn from the distribution.
- ▶ The samples drawn from the distribution are called particles, and are denoted by:

$$X_t := x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[M]}$$

Each particle is a hypothesis as to what the true world state may be at time  $t$ .

$$x_t^{[m]} \sim p(x_t | z_{1:t}, u_{1:t})$$

As a consequence, as  $M \rightarrow \infty$ , the denser a sub-region of the state space is populated by samples, the more likely it is that the true state falls into this region.

- As with the previous algorithms, the particle filter algorithm estimates the belief  $bel(x_t)$  recursively from the belief  $bel(x_{t-1})$ .

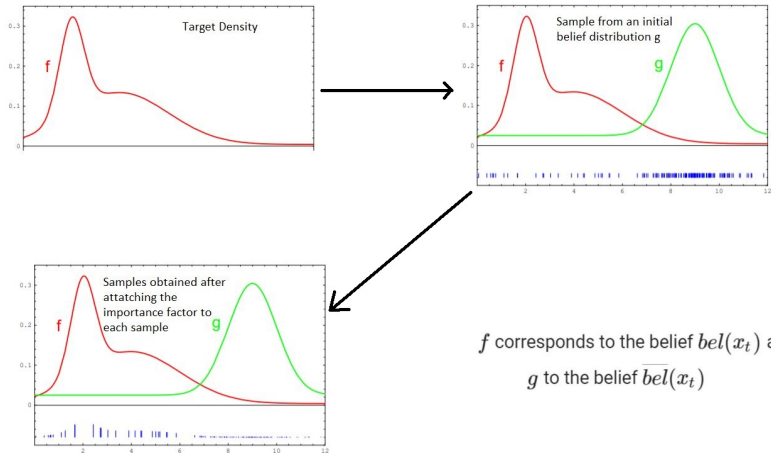
```
1:  Algorithm Particle_filter( $\mathcal{X}_{t-1}, u_t, z_t$ ):  
2:     $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset$   
3:    for  $m = 1$  to  $M$  do  
4:      sample  $x_t^{[m]} \sim p(x_t \mid u_t, x_{t-1}^{[m]})$   
5:       $w_t^{[m]} = p(z_t \mid x_t^{[m]})$   
6:       $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t + \langle x_t^{[m]}, w_t^{[m]} \rangle$   
7:    endfor  
8:    for  $m = 1$  to  $M$  do  
9:      draw  $i$  with probability  $\propto w_t^{[i]}$   
10:     add  $x_t^{[i]}$  to  $\mathcal{X}_t$   
11:    endfor  
12:    return  $\mathcal{X}_t$ 
```

- ▶ The set of particles resulting from iterating Step 4  $M$  times is the filter's representation of  $\overline{bel}(x_t)$
- ▶ Line 5 calculates for each particle  $x_t^{[m]}$  the so-called importance factor, denoted by  $w_t^{[m]}$ . Importance factors are used to incorporate the measurement  $z_t$  into the particle set.



- ▶ lines 8-11 implement what is known as *resampling* or importance resampling.
  - The algorithm draws with replacement  $M$  particles from the temporary set  $\overline{X}_t$ . The probability of drawing each particle is given by its importance weight.
  - Resampling transforms a particle set of  $M$  particles into another particle set of the same size.
  - By incorporating the importance weights into the resampling process, the distribution of the particles change: whereas before the resampling step, they were distributed according to  $\overline{bel}(x_t)$ , after the resampling they are distributed (approximately) according to the posterior
$$bel(x_t) = \eta p(z_t | x_t^{[m]}) \overline{bel}(x_t)$$

## An Illustration of the Particle Filter Algorithm



$f$  corresponds to the belief  $bel(x_t)$  and  $g$  to the belief  $\bar{bel}(x_t)$

### The Binary Bayes Filter (with static state)

- ▶ Certain problems are best formulated as binary state problems (i.e with states  $x$  and  $\neg x$ ), where the robot needs to estimate a static state from a sequence of sensor measurements.
- ▶ Since the state is static, the belief is a function of the measurements

$$bel_t(x) = p(x|z_{1:t}, u_{1:t}) = p(x|z_{1:t})$$

Note, in such problems:

$$bel_t(\neg x) = 1 - bel_t(x)$$

The lack of time index indicates static state.

- ▶ The belief in such problems is commonly implemented as a log odds ratio.

Where, odds of an event  $x$  is defined as:  $\frac{p(x)}{p(\neg x)} = \frac{p(x)}{1-p(x)}$

log odds of this expression is:


$$l(x) := \log \frac{p(x)}{1-p(x)}$$

- ▶ The Bayes filter for updating beliefs in log odds representation is computationally elegant.

```
1: Algorithm binary_Bayes_filter( $l_{t-1}, z_t$ ):  
2:    $l_t = l_{t-1} + \log \frac{p(x|z_t)}{1-p(x|z_t)} - \log \frac{p(x)}{1-p(x)}$   
3:   return  $l_t$ 
```

## Non-Parametric Filters

- ▶ This binary Bayes filter uses an inverse measurement model  $p(x|z_t)$ , instead of the familiar forward model  $p(z_t|x)$ .
- ▶ Inverse models are often used in situations where measurements are more complex than the binary state.
- ▶ For example, it is easier to devise a function that calculates a probability of a door being closed from a camera image, than describing the distribution over all camera images that show a closed door. In other words, it is easier to implement an inverse than a forward sensor.



- ▶ Here the state is extremely simple, but the space of all measurements is huge.

The belief can be recovered from the log odds ratio by the following equation:

$$bel_t(x) = 1 - \frac{1}{1 + \exp(l_t(x))}$$

### Correctness of the Binary Bayes Filter Algorithm

$$bel(x) = p(x|z_{1:t}, u_{1:t}) = p(x|z_{1:t})$$

$$\text{Using Bayes rule } \{p(A|B) = \frac{p(B|A)p(A)}{p(B)}\},$$

$$= \frac{p(z_t|x, z_{1:t-1})p(x|z_{1:t-1})}{p(z_t|z_{1:t-1})}$$

$$= \frac{p(z_t|x)p(x|z_{1:t-1})}{p(z_t|z_{1:t-1})}$$

Applying Bayes rule to the measurement model,

$$p(z_t|x) = \frac{p(x|z_t)p(z_t)}{p(x)}$$

Substituting,

$$p(x|z_{1:t}) = \frac{p(x|z_t)p(z_t)p(x|z_{1:t-1})}{p(x)p(z_t|z_{1:t-1})}$$

Similarly,

$$p(\neg x|z_{1:t}) = \frac{p(\neg x|z_t)p(z_t)p(\neg x|z_{1:t-1})}{p(\neg x)p(z_t|z_{1:t-1})}$$

---

Consequently,

$$\begin{aligned}\text{odds} &= \frac{p(x|z_{1:t})}{p(\neg x|z_{1:t})} = \frac{p(x|z_t)}{p(\neg x|z_t)} \frac{p(x|z_{1:t-1})}{p(\neg x|z_{1:t-1})} \frac{p(\neg x)}{p(x)} \\ &= \frac{p(x|z_t)}{1-p(x|z_t)} \frac{p(x|z_{1:t-1})}{1-p(x|z_{1:t-1})} \frac{1-p(x)}{p(x)}\end{aligned}$$

---

$$\begin{aligned}\text{So, } l_t(x) &= \log \frac{p(x|z_t)}{1-p(x|z_t)} + \log \frac{p(x|z_{1:t-1})}{1-p(x|z_{1:t-1})} + \log \frac{1-p(x)}{p(x)} \\ &= \log \frac{p(x|z_t)}{1-p(x|z_t)} - \log \frac{p(x)}{1-p(x)} + l_{t-1}(x)\end{aligned}$$